

This material takes 1 hour.

1 Persistent Data Structures

Sarnak and Tarjan, "Planar Point Location using persistent trees", Communications of the ACM 29 (1986) 669–679

"Making Data Structures Persistent" by Driscoll, Sarnak, Sleator and Tarjan
Journal of Computer and System Sciences 38(1) 1989

Idea: be able to query and/or modify past versions of data structure.

- ephemeral: changes to struct destroy all past info
- partial persistence: changes to most recent version, query to all past versions
- full persistence: queries and changes to all past versions (creates "multiple worlds" situation)

Goal: general technique that can be applied to *any* data structure.

Application: planar point location.

- a computational geometry foreshadow
- planar subdivision
 - division of plane into polygons
 - query: "what polygon contains this point"
 - define complexity of input?
 - edges of polygons form n segments/edges
 - segments meet only at ends/vertices
- numerous special-purpose solutions
- One solution: dimension reduction
 - vertical line through each vertex
 - divides into slabs
 - in slab, segments maintain one vertical ordering
 - find query point slab by binary search
 - build binary search tree for slab with "above-below" queries
 - n binary search trees, size $O(n^2)$, time $O(n^2 \log n)$
- observation: trees all very similar
- think of x axis as time, slabs as "epochs"
- at end of epoch, "delete" segments that end, "insert" those that start.

- over all time, only n inserts, n deletes.
- must be able to query over all times

Persistent sorted sets:

- $\text{find}(x, s, t)$ find (largest key below) x in set s at time t
- $\text{insert}(i, s, t)$ insert i in s at time t
- $\text{delete}(i, s, t)$.

We use partial persistence: updates only in “present”

Implement via persistent search trees.

Result: $O(n)$ space, $O(\log n)$ query time for planar point location.

2 Persistent Trees

Many data structures

- consist of fixed-size *nodes* connected by *pointers*
- accessed from some *root* entry point
- data structure ops traverse pointers
- and modify node fields/pointers
- runtime determined by number of traverse/modify steps

Our goal: use *simulation* to transform any ephemeral data structure into a persistent one

- measure of success: time and space cost for each *primitive traversal* and *primitive modification* step.
- An exposed data structure operation will do some number of primitive steps
- We'll limit to trees.
- But ideas generalize

Full copy bad.

Fat nodes method:

- replace each (single-valued) field of data structure by list of all values taken, sorted by time.
- requires $O(1)$ space per data change (**unavoidable** if keep old data)
- to lookup data field, need to search based on time.
- store values in binary tree

- checking/changing a data field takes $O(\log m)$ time after m updates
- **multiplicative** slowdown of $O(\log m)$ in structure access.

Path copying:

- can only reach node by traversing pointers starting from root
- changes to a node only visible to *ancestors* in pointer structure
- when change a node, copy it and ancestors (back to root of data structure)
- keep list of roots sorted by update time
- $O(\log m)$ time to find right root (or const, if time is integers) (**additive** slowdown)
- same access time as original structure
- *additive* instead of multiplicative $O(\log m)$.
- modification time and space usage equals number of ancestors: possibly huge!

Combined Solution (trees only):

- fat nodes bad, but “pleasingly plump” ok.
- use fat nodes, but when get too fat, make new path copy
- in each node, store 1 *extra* time-stamped field
- if full, overrides one of standard fields for any accesses later than stamped time.
- access rule
 - standard access, just check for overrides while following pointers
 - constant factor increase in access time.
- update rule:
 - when need to change/copy pointer, use extra field if available.
 - otherwise, make new copy of node with new info, and recursively modify parent.
- Analysis
 - *live* node: pointed at by *current* root.
 - potential function: number of *full* live nodes.
 - copying a node is free (new copy not full, pays for copy space/time)

- pay for filling an extra pointer (do only once, since can stop at that point).
- amortized space per update: $O(1)$.

Power of twos: Like Fib heaps. Show binary tree of modifications.

Application: persistent red-black trees:

- aggressive rebalancers
- amortized cost $O(1)$ to change a field.
- store red/black bit in each node
- use red/black bit to rebalance.
- depth $O(\log n)$
- search: standard binary tree search; no changes
- update: causes changes in red/black fields on path to item, $O(1)$ rotations.
- result: $(\log n)$ space *per insert/delete*
- geometry does $O(n)$ changes, so $O(n \log n)$ space.
- $O(\log n)$ query time.

Improvement:

- red-black bits used only for updates
- only need current version of red-black bits
- don't store old versions: just overwrite
- only persistent updates needed are for $O(1)$ rotations
- so $O(1)$ space per update
- so $O(n)$ space overall.

Result: $O(n)$ space, $O(\log n)$ query time for planar point location.

Extensions:

- method extends to arbitrary pointer-based structures.
- $O(1)$ cost per update for any pointer-based structure with any constant indegree.
- full persistence with same bounds.