

1 Online algorithms

Motivation:

- till now, our algorithms start with input, work with it
- (exception: data structures—come back later)
- now, suppose input arrives a little at a time, need instant response
- eg stock market, paging
- question: what is a “good” algorithm.
- depends on what we measure.
- if knew whole input σ in advance, easy to optimize $C_{MIN}(\sigma)$
- ski rental problem: rent 1, buy T . don’t know how often we are going to ski.
- notice that on some inputs, can’t do well! (stock market that only goes down, thrashing in paging)
- problem isn’t to decide fast, rather what to decide.

Definition: competitive ratio

- compare to full knowledge optimum
- k -competitive if for all sequences etc. $C_A(\sigma) \leq kC_{MIN}(\sigma)$
- sometimes, to ignore edge effects, $C_A(\sigma) \leq kC_{MIN}(\sigma) + O(1)$.
- idea: “regret ratio”
- analyze ski rental
 - Only choice is after how many days to buy.
 - If I rent d days and then buy (total cost of $d + T$) what is worst competitive ratio for you?
 - Before I buy skis, ratio is 1 if adversary knows should be renting, and keeps getting worse if adversary has bought. Either way, ratio not improving.
 - Once I have bought, my cost doesn’t increase and adversary’s doesn’t decrease, so ratio doesn’t worsen
 - combine: worst ratio is at moment I have bought skis.
 - I pay $d + T$, adversary pays $\min(d, T)$
 - Ratio $(d + T) / \min(d, T)$ optimized at $d = T$.

- we think of competitive analysis as a (zero sum) game between algorithm and adversary. Want to find the best strategy for algorithm.
- supposed to be competitive against all sequences. So, can imagine that adversary is adapting to algorithm's choices (to get worst sequence)

1.1 Finance

Known or unknown duration. But assume know which offer is last.
Need fluctuation ratio ϕ between largest M and smallest m price.
Selling peanuts:

- Break into $\log \phi$ groups of equal amounts
- Sell group i for value $m \cdot 2^i$
- One group sold for at least half of max price
- So achieve $\log \phi$ competitive

Selling (one) car: Best deterministic algorithm: agree to first price exceeding \sqrt{Mm}

- $\sqrt{\phi}$ competitive
- note have to know when last offer

Can achieve $\log \phi$ randomized

- Consider powers of 2 between m and M
- Choose one at random
- sell all at first bid exceeding
- with prob $1/\log \phi$, pick the power of 2 that is within factor 2 of highest offered price.
- even if know ϕ but don't know m , can just run above alg after seeing first price

Graham's Rule

Define $P|| \max C_j$ to minimize load.

NP-complete to solve exactly!

Always assign to least loaded machine:

- any alg has 2 lower bounds: average load and maximum job size.
- Suppose M_1 has max load L , let p_j be biggest job.
- claim every machine has $L - p_j$ (else wouldn't have assigned last job to M_1)

- thus total load at least $\sum p_i = m(L - p_j) + p_j$
- thus $\text{OPT} \geq L - p_j + p_j/m$
- but $\text{OPT} \geq p_j$, so $(2 - 1/m)\text{OPT} \geq L$

More recent algs do somewhat better:

- keep some machines small
- algorithms not too bad, proofs awful!
- Bartal et al '95: $2 - 1/70 = 1.986$
- Karger et al '96: 1.945
- Albers '97: 1.923

1.2 Move to front

Studying heuristics for reorganizing a list after you access it

- a natural heuristic: move accessed item to front
- is it a good heuristic?
- compare to omniscient algorithms that can move accessed item anywhere
- so e.g. *won't* move item to front if not accessed again soon

Potential function: number of inversions.

- amortized cost
- suppose search for item x_j at j in opt, at k in MTF
- suppose v items precede in MTF but not OPT
- then $k - v - 1$ precede in BOTH
- so $k - v - 1 \leq j - 1$ so $k - v \leq j$
- MTF creates $k - v - 1$ new inversions and kills v old ones,
- so amortized cost is $k + (k - v - 1) - v \leq 2(k - v) \leq 2j$
- now do opt's move.
- moving x_j towards front only decreases inversions (since x_j already at front in MTF)
- so cannot increase potential
- so doesn't increase amortized cost of access

Strengthen

- we can allow adversary to make *arbitrary transposes*
- if we charge the adversary 1 for each transpose
- since then any increase in potential (and thus our amortized cost) is upper bounded by the adversary's cost
- so we remain 2-competitive

Lower bound:

- suppose n items in list
- nasty algorithm: always request last in list
- generates a sequence of length m
- total cost mn
- consider alg that picks random order and doesn't change
- expected access time for each item is $(n - 1)/2$
- expected total cost $m(n - 1)/2$
- some algorithm achieves that cost
- offline alg can use that one
- ratio $2n/(n - 1) \rightarrow 2$

Note: our strengthened bound assumes opt is paying 1 for transposes

- if opt can rearrange anything it sees, can't beat $\Omega(n/\log n)$ competitive
- "Order by Next Request" heuristic—rearrange everything you pass
- Munro 2000
- achieves cost equal to entropy
- so consider a uniform random sequence
- entropy is $O(\log n)$
- an online algorithm will pay $n/2$ in expectation every time.
- so $\Omega(n/\log n)$ competitive

1.3 Paging problem

- define
- LRU, FIFO, LIFO, Flush when full, Least Freq Use
- LIFO, LFU not competitive
- LRU, FIFO k -competitive.
- will see this is best possible (det)

LRU is $k/(k - h + 1)$ competitive against h -page adversary

- k -phase partition: maximum groups containing k pages
- ie, next request is $(k + 1)$ st page.
- at most k faults by LRU (or FIFO, or flush when full)
- Let q be first page of phase i
- In rest of phase, up to and including first request of phase i , OPT has $h - 1$ pages not including q
- so must have $k - h + 1$ faults
- ratio $k/(k - h + 1)$.

Observations:

- proved without knowing optimum
- instead, derived *lower bound* on cost of *any* algorithm
- same argument applies to FIFO.

Lower bound: no online algorithm beats k -competitive.

- set of $k + 1$ pages
- always ask for the one A doesn't have
- faults every time.
- so, just need to show can get away with 1 fault every k steps
- have k pages, in memory. When fault, look ahead, one of $k + 1$ isn't used in next k , so evict it.
- one fault every k steps
- so A is only k -competitive.

Observations:

- Lower Bound can be proven without knowing OPT, often is.
- competitive analysis doesn't distinguish LRU and FIFO, even though know different in practice.
- still trying to refine competitive analysis to measure better: new SODA paper: "LRU is better than FIFO"
- applies even if just have $k + 1$ pages!

Optimal offline algorithm: Longest Forward Distance

- evict page that will be asked for farthest in future.
- suppose MIN is better than LFD. Will make NEW, as good, agrees more with LFD.
- Let σ_i be first divergence of MIN and LFD (at page fault)
- LFD discards q , MIN discards p (so p will be accessed before q after time i)
- Let t be time MIN discards q
- revise schedule so MIN and LFD agree up to t , yielding NEW
- NEW discards q at i , like LFD
- so MIN and NEW share $k - 1$ pages. will preserve till merge
- in fact, q is unique page that MIN has that NEW doesn't
- case 1: $\sigma_i, \dots, \sigma_t, \dots, p, \dots, q$
 - until reach q
 - let e be unique page NEW has that MIN doesn't (init $e = p$)
 - when get $\sigma_l \neq e$, evict same page from both
 - note $\sigma_l \neq q$, so MIN does fault when NEW does
 - both fault, and preserves invariant
 - when $\sigma_l = e$, only MIN faults
 - when get to q , both fault, but NEW evicts e and converges to MIN.
 - clearly, NEW no worse than MIN
- case 2: t after q
 - follow same approach as above till hit q
 - since MIN didn't discard q yet, it doesn't fault at q , but
 - since p requested before q , had $\sigma_l = e$ at least once, so MIN did *worse* than NEW. (MIN doesn't have p till faults)

- so, fault for NEW already paid for
- still same.
- prove that can get to LFD without getting worse.
- so LFD is optimal.

2011 End lecture 20

Randomized Online Algorithms

An online algorithm is a two-player zero sum game between algorithm and adversary. Well known that optimal strategies require randomization.

A *randomized online algorithm* is a probability distribution over deterministic online algorithms.

- idea: if adversary doesn't know what you are doing, can't mess you up.
- idea: can't see adversary's "traps", but have certain probability of wiggling out of them.
- in practice, don't randomly pick 1 det algorithm at start. Instead, make random choices on the way. But retrospectively, gives 1 deterministic algorithm.

Algorithm is k -competitive if for any σ , $E[C_A(\sigma)] \leq k \cdot OPT + O(1)$.

Adversaries:

- **oblivious:** knows probability distribution but not coin tosses. Might as well pick input in advance.
- **fully adaptive:** knows all coin tosses. So algorithm is deterministic for it.
- **adaptive:** knows coin tosses up to present—picks sequence based on what did.
- clearly adaptive stronger than oblivious.
- oblivious adversary plausible in many cases (eg paging)
- problematic if online behavior affects nature (eg, paging an alg that changes behavior if it sees itself thrashing)
- for now, oblivious

1.3.1 Randomized Paging

Idea: evict random page?

- k -competitive against *adaptive* adversary
- but uses no memory
- trading space for randomness

Marking algorithm:

- initially, all pages marked (technicality)
- on fault, if all marked, unmark all
- evict random unmarked page
- mark new page.

Fiat proved: Marking is $O(\log k)$ competitive for k pages.

Phases:

- first starts on first fault
- ends when get $k + 1^{st}$ distinct page request.
- so a phase has k distinct pages
- cost of M is cost of phases
- note: defined by input, independent of coin tosses by M
- but, marking tracks:
 - by induction, unmark iff at end of phase
 - by induction, all pages requested in phase stay marked till end of phase
 - so, pay for page (if at all) only on first request in phase.
 - by induction, at end of phase memory contains the k pages requested during the phase.

Analysis:

- ignore all but first request to a page (doesn't affect M , helps offline)
- compare phase-by-phase cost
- phase i starts with S_i (ends with S_{i+1})
- request *clean* if no in S_i . M must fault, but show offline pays too
- request *stale* if in S_i . M faults if evicted during phase. Show unlikely.

Online cost:

- Expected cost of stale request:
 - suppose had s stale and c clean requests so far.
 - so s pages of S_i known to be currently in memory
 - remaining $k - s$ may or may not be.
 - in particular, c of them got evicted for clean requests
 - what prob current request was evicted? $c/(k - s)$
 - this is expected cost of stale request.

- Cost of phase.
 - Suppose has c_i clean requests, $k - c_i$ stale.
 - Pay c_i for clean.
 - for stale requests, pay at most

$$c_i \left(\frac{1}{k} + \frac{1}{k-1} + \cdots + \frac{1}{c_i+1} \right) = c_i (H_k - H_{c_i})$$

- so total at most $c_i \log k$

Offline cost:

- potential function Φ_i = difference between M and O (offline) at start of phase i .
- got c_i clean requests, not in M 's memory. So at least $c_i - \Phi_i$ not in O 's memory.
- at end of round, M has all k most recent requests. So O is missing Φ_{i+1} of k this round's requests. Must have evicted (thus paid for) them.
- so, $C_i(O) \geq \max(c_i - \Phi_i, \Phi_{i+1}) \geq \frac{1}{2}(c_i + \Phi_i - \Phi_{i+1})$.
- sum over all phases; telescopes. Deduce $C_i \geq \frac{1}{2} \sum c_i$.

Summary: If online pays $x \log k$, offline pays $x/2$. So, $(\log k)$ -competitive.

Lower Bounds

Turns out that $O(\log k)$ is tight for randomized algorithms (Fiat). How prove? Recall that situation is a game:

- in general, optimal strategy of both sides is randomized
- online chooses random alg, adversary chooses random input
- leads to payoff matrix—expected value of game

- number in matrix is cost for alg on that input
- Von Neumann proved equality of minimax and maximins
- notice: player who picks strategy second can use deterministic choice
- note when one player's strategy known, other player can play deterministically to meet optimum.
- above, assumed adversary knew online's strategy, so he played deterministically
- for lower bound, we let adversary have randomized strategy, look for best deterministic counter!
- If give random input for which no deterministic alg does well, we get a lower bound.

Formalize:

- say online A is c -competitive against an input distribution p_σ if $E_\sigma(C_A(\sigma)) \leq cE_\sigma(C_{OPT}(\sigma))$ (note: OPT gets to see sequence before going)
- Theorem: if for some distribution no deterministic alg is c -competitive, then no randomized algorithm is c -competitive.
- to prove, suppose have c -competitive randomized alg, show det c -competitive against any σ .
- consider payoff $E_A[C_A(\sigma) - cC_{OPT}(\sigma)]$
- by assumption, some dist on A achieves non-positive payoff.
- remains true even if choose best possible randomized strategy on σ
- once do so, have deterministic counter A
- so for any p_σ on σ , some A such $E_\sigma[C_A(\sigma) - cC_{OPT}(\sigma)] \leq 0$
- in other words, A is c -competitive against p_σ .

For paging:

- set of $k + 1$ pages
- uniform random sequence of requests
- *any* deterministic (or randomized!) algorithm has an expected $1/k$ fault per request. So cost n/k if seq length n
- what is offline cost? on fault, look ahead to page that is farthest in future.
- *phase* ends when all $k + 1$ pages requested

- offline faults once per phase
- how long is a phase? coupon collection. $\Omega(k \log k)$.
- intuitively, number of faults is $n/k \log k$
- formally, use “renewal theory” that works because phase lengths are i.i.d.
- deduce expected faults $n/k \log k$, while online is n/k
- $\log k$ gap, so online not $\log k$ -competitive.

2011 End Lecture 21

1.4 k -server

Definition:

- metric space with k servers on points
- request is a point in space
- must move a server, cost is distance.
- eg taxi company
- paging is special case: all distances 1, servers on “memory pages”
- also multi-head disks
- compute offline by dynamic program or reduction to min cost flow

Greedy doesn’t work:

- 2 servers, 1 far away, other flips between 2 points.
- need an algorithm that moves a far away server sometimes in case a certain region is popular

Fancy algorithmics:

- HARMONIC: randomized, move with probability inversely proportional to distance from goal
- WORK FUNCTION: track what offline algorithms would have done (computationally very expensive) and then do your best to move into a similar configuration.
- in 2001, work-function was proven $2k$ -competitive using a black magic potential function
- conjectured k -competitive.
- questions remain on finding an algorithm that does little work per input.

1.4.1 On a Line

greedy algorithm bad if requests alternate a near b but server on distant c .
double coverage algorithm (DC):

- If request outside convex hull, move nearest point to it.
- else, move nearest point on each side towards it equal distance till one hits.

k -competitive.

- let M be min-cost matching of opt points to DC points
- $\Phi = kM + \sum_{i < j} d(s_i, s_j)$
- show:
 - adversary moves d : increases Φ by $\leq kd$
 - DC moves moves d : decrease Φ by d
- deduce: DC is k -competitive because it moves only k times opt.

Analysis:

- adv moves d just increases M by d , so $\Delta\Phi \leq kd$
- DC moves d .
- If to outside hull, note adversary already has a point at destination; moving point must match to it (else matches something else; uncross).
- so $\Delta M = -d$ while $\delta\Sigma = (k-1)d$. claim follows: $\Delta\phi = -kd + (k-1)d = -d$
- if inside hull, one of moving points is matched to request. So that move decreases M . Other move may increase M at most by same amount, so no change to M .
- Now consider Σ . Moves of two points cancel out with respect to other points, but they get $2d$ units closer.